

Performance Co-Pilot cheat sheet

<https://github.com/christianhorn/pcp-cheat-sheet>

PCP basics

Installation

Package **zero-conf** pulls in dependencies, starts daemons, starts archiving of a default set of metrics. On RPM based distros (RHEL, Fedora, CentOS etc.):

```
# dnf -y install pcp-zeroconf
```

Installation via Ansible playbook:

```
# linux-system-roles.metrics
```

Verify pcp installation:

```
# pcp
# systemctl status pcpd pmlogger
```

Important tools

Tools from package 'pcp-system-tools' can be used with either the running pcpd, or PCP archive files:

- pcp atop
- pcp collectl
- pcp free
- pcp iostat
- pcp dstat

Important Pathes

```
/var/lib/pcp/config/
/var/log/pcp/
/etc/pcp/
```

Working with metrics

Which metrics are offered by the running pcpd?

```
# pminfo
```

Which metrics related to cpu are available?

```
# pminfo | grep cpu
```

pmie

pmie, performance metrics interference engine, can react on defined metric states: send email on high load, and so on.

```
# pmie --verbose --timestamp --interval 1
# /etc/pcp/pmie/config.default
# pmie --archive 20200512 --config <rules>
```

pmdas

PMDA installation

PMDA's are code pieces capable of reading metrics from their area like sensors, database, and so on. PMDA's can be searched as packages and installed, for example on yum4/dnf distros:

```
# dnf search pcp-pmda
# dnf install -y pcp-pmda-lmsensors
# cd /var/lib/pcp/pmdas/lmsensors
# less README
# ./Install
```

Anomaly search in archives

Which metrics are remarkably different in a certain timeframe? Example: we had I/O problems from 2am to 3am:

```
# cat cull
rsyslog.
# pmdiff -z -X ./cull -q 10 \
--start @10:00 --finish @10:30 \
--begin @12:00 --end @12:30 \
./archives/20120512 | less
```

What are the top 5 cpu and memory hogs?

```
# pmrep proc.hog.cpu, proc.memory.rss \
-J 5 -1 -g -b MB
```

What are the current PIDs, and how much rss uses process with pid 75?

```
# pmrep proc.smaps.rss -g | less
# pmrep proc.smaps.rss -g -i '*.75.*'
```

Remote collection

Install PCP on clients

Setup client systems to offer metrics via pcpd: install pcp, open packet based firewall, enable remote access in pcpd:

```
# yum -y install pcp
# firewall-cmd --permanent --zone=public \
--add-port=44321/tcp
# firewall-cmd --reload
# CONFIG=/etc/sysconfig/pcpd
# if grep -q PMCD_LOCAL $CONFIG; then \
sed -ie 's,PMCD_LOCAL.*,PMCD_LOCAL=0,' \
/etc/sysconfig/pcpd \
else
echo 'PMCD_LOCAL=0' »/etc/sysconfig/pcpd
fi
# grep PMCD_LOCAL /etc/sysconfig/pcpd
# service pcpd restart
# chkconfig pcpd on
```

Install PCP on collector system

On the collector, we install pcp-zeroconf which also sets up logging to archive files. We then set variable CLIENT to the clients name, create a config- and controlfile, and notify pmlogger of the changes.

```
# yum -y install pcp-zeroconf
# CLIENT=rhel7u8a
# /usr/libexec/pcp/bin/pmlogconf \
/var/lib/pcp/config/pmlogger/config.$CLIENT
# echo "$CLIENT.local n n " \
PCP_LOG_DIR/pmlogger/$CLIENT.local" \
-r -T30d -c config.$CLIENT" \
>/etc/pcp/pmlogger/control.d/$CLIENT
# /usr/libexec/pcp/bin/pmlogger.check
```

Live metrics

Accessing kernel.all.load metric of local system:

```
# pmrep -p kernel.all.load
```

Get live metrics from 2 remote systems:

```
# pmstat -h local: -h remotesystem1
```

Links

- [1] kbase: Index of (PCP) articles, solutions, tutorials, white papers
- [2] Ansible:<https://github.com/linux-system-roles/metrics> <https://github.com/performancecopilot/ansible-pcp>
- [3] Performance Co-Pilot site
- [4] Articles: Solve performance mysteries with PCP / PCP and podman / PCP and dstat

Archive files

Basics

Which archive is pmlogger logging into?

```
# pcp
```

Set a variable to current archive, and evaluate how many metrics are logged in the archive:

```
# cd /var/log/pcp/pmlogger/<hostname>
# pminfo -a <archivename> | wc -l
```

Have pmdiff point out 'significant peaks' in archives:

```
# pmdiff -a <archivename>
```

Merge all archive files ending in .0 into a single archive:

```
# pmlogextract *.0 single-archive-name
```

Accessing metrics

Which time span is covered by the archive file, time displayed for the timezone you are running while analyzing?

```
# pmdumplog -L <archive>
```

Which time span is covered by the archive file, time displayed

with the timezone of the system which recorded the archive file?

```
# pmdumplog -L <archive>-z
```

Most basic access to metrics:

```
# pmstat -t 1
# pmstat -a <archivename><metric>
# pmrep -a <archivename><metric>
# pcp -a 20180831.11.31 --origin @1pm \
dstat --time --disk --mem 60sec 10
# pmrep -a 20220128 kernel.cpu -p -t 300 \
--hostzone -S @19:30
```

Show I/O stats for sda, starting at Jul 19th

```
# pmiostat -z -a <archive>-x t \
-t 10m -S'@Jul 19 19:55:00' \
-s 10 -P0 -R 'sda$'
```

Graphical access

```
# pmchart
# dnf -y install pcp2pdf; \
pcp2pdf -a <archivename>
```